


---

DuyHai DOAN, Technical Advocate

# Shameless self-promotion

Duy Hai DOAN

- , ,
- - (Achilles, ) :
- 
-  **duy\_hai.doan@datastax.com**
- production

•

# Agenda

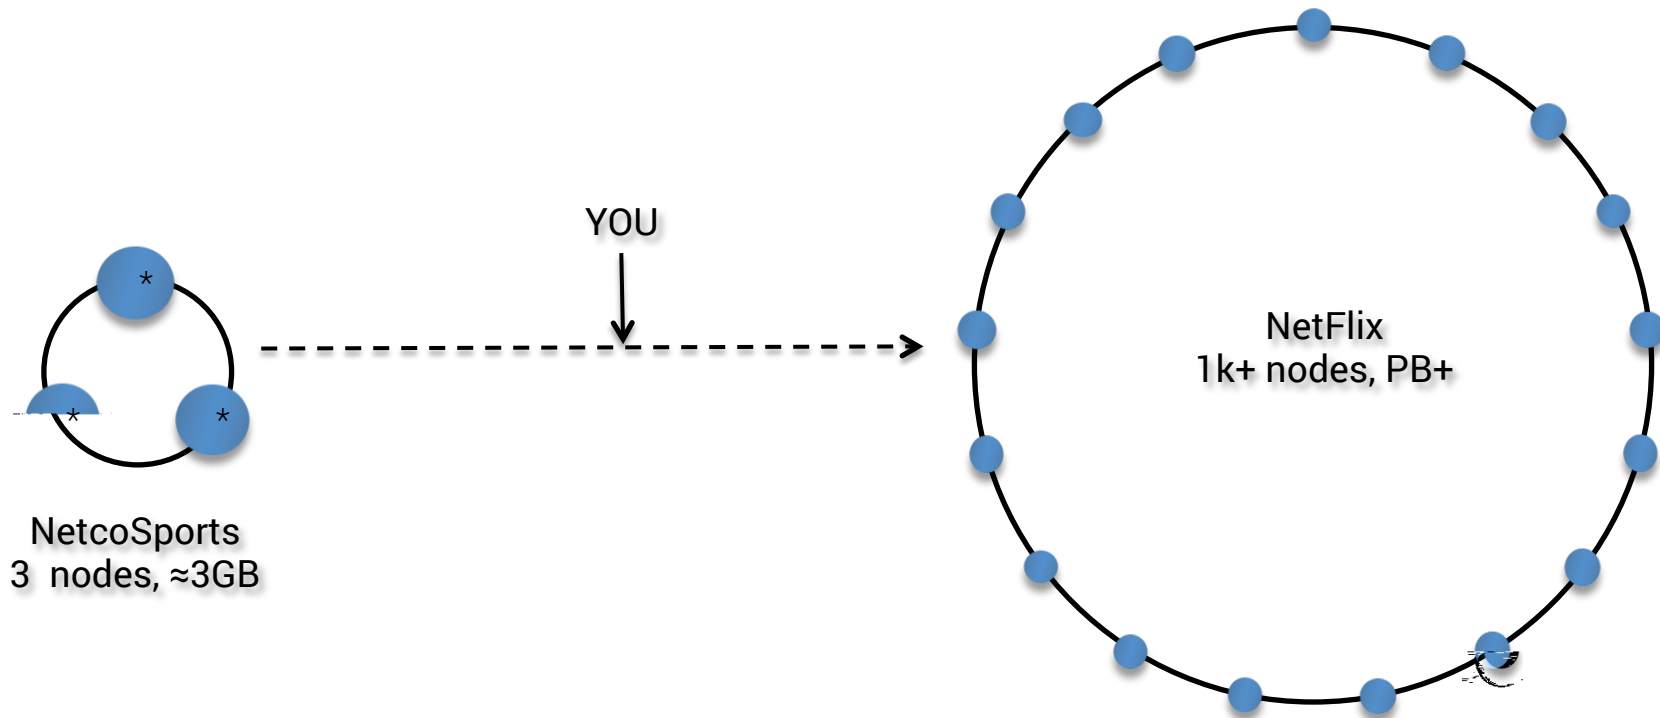
## Architecture

- 
- 

## Data model

- ( ),
- ( , , , )
-

# Cassandra 5 key facts



# Cassandra 5 key facts

- $(\approx 100\% - )$   
( )

# Cassandra 5 key facts

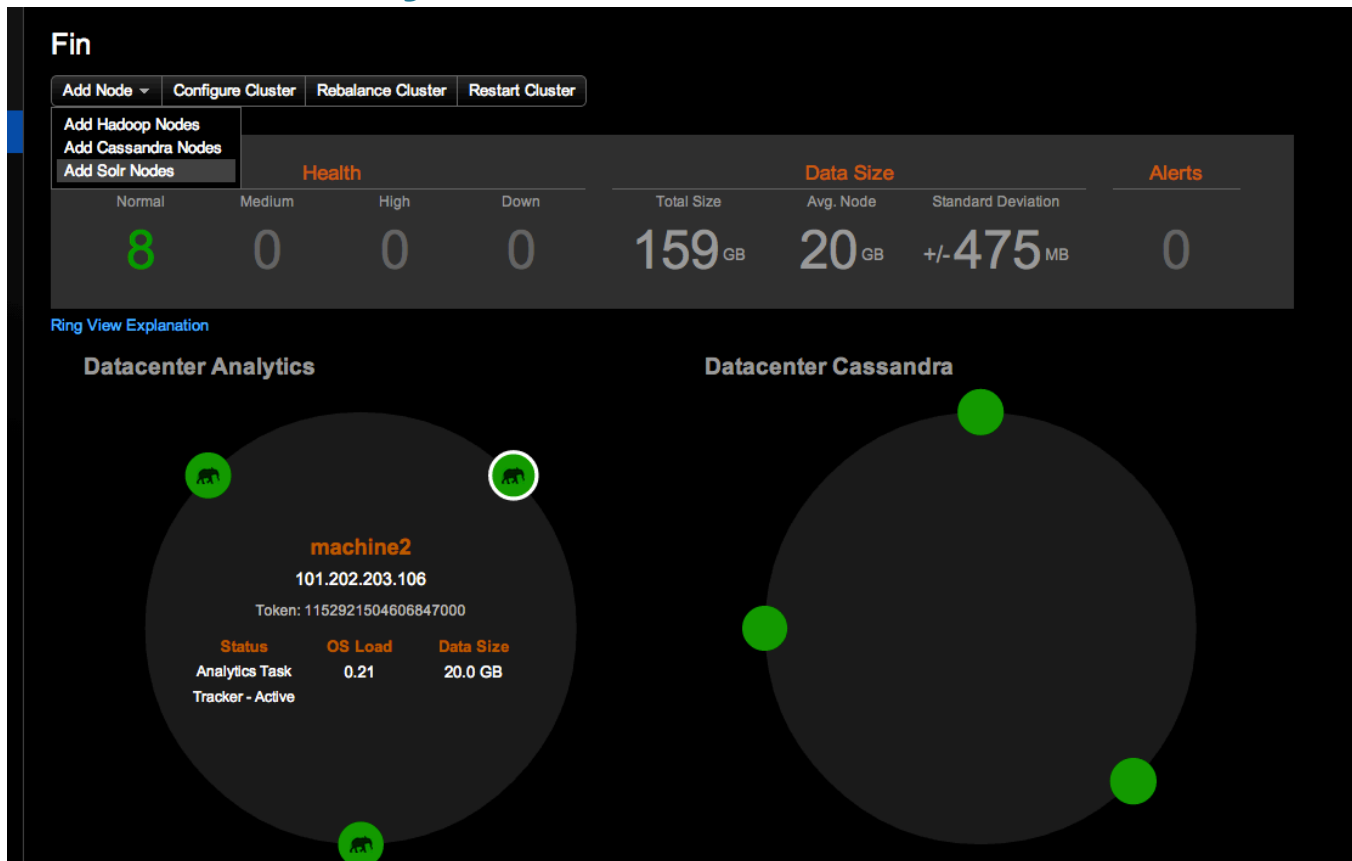
- 
- - - - ( )
- -
- /

# Cassandra 5 key facts

- 1 = 1 + 1
- 
-



# Cassandra 5 key facts



# Cassandra 5 key facts

- Cassandra + Spark = awesome !
- realtime streaming

# Cassandra architecture

Cluster  
Replication

# Cassandra architecture

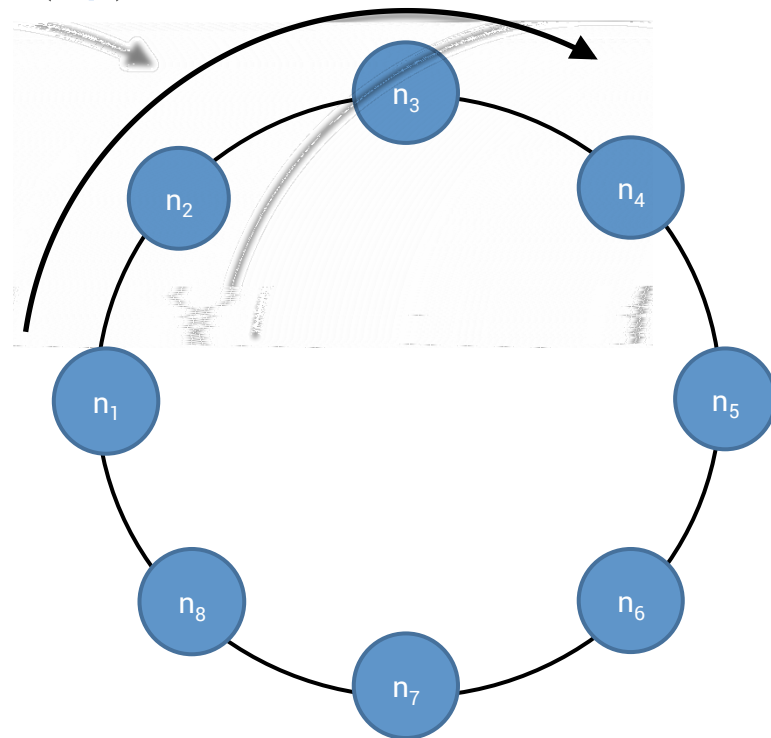
- DynamoDB
- masterless
  -
- Big Table
- /

# Data distribution

: #partition  $\rightarrow$  token =  $(\#p)$

: - ,

=  $(2^{64}/2)$



# Token Ranges

A: 0, /8

B: /8, 2 /8

C: 2 /8, 3 /8

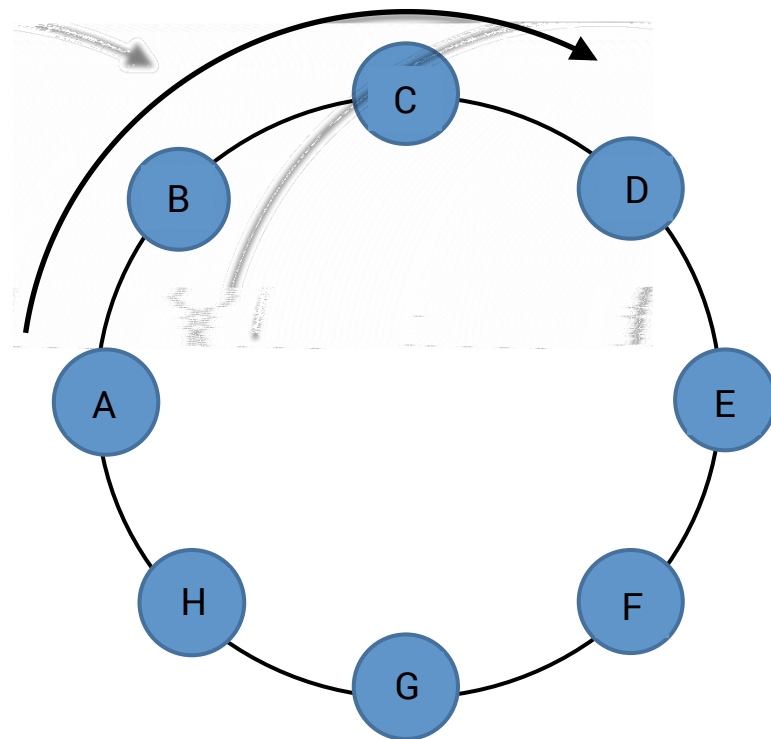
D: 3 /8, 4 /8

E: 4 /8, 5 /8

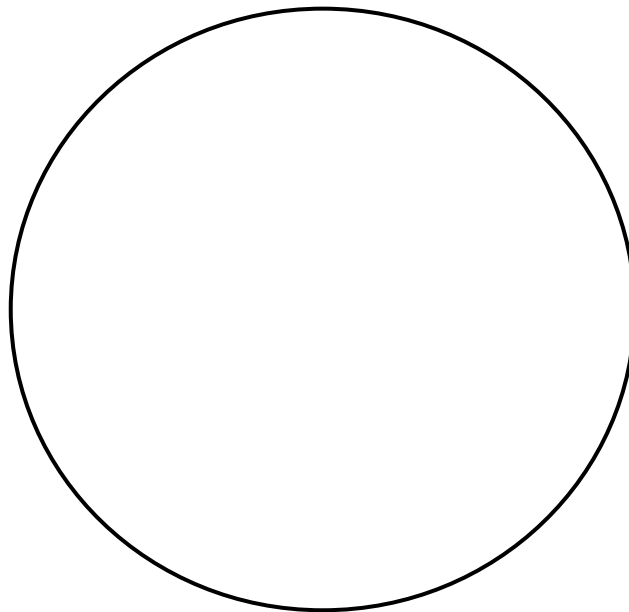
F: 5 /8, 6 /8

G: 6 /8, 7 /8

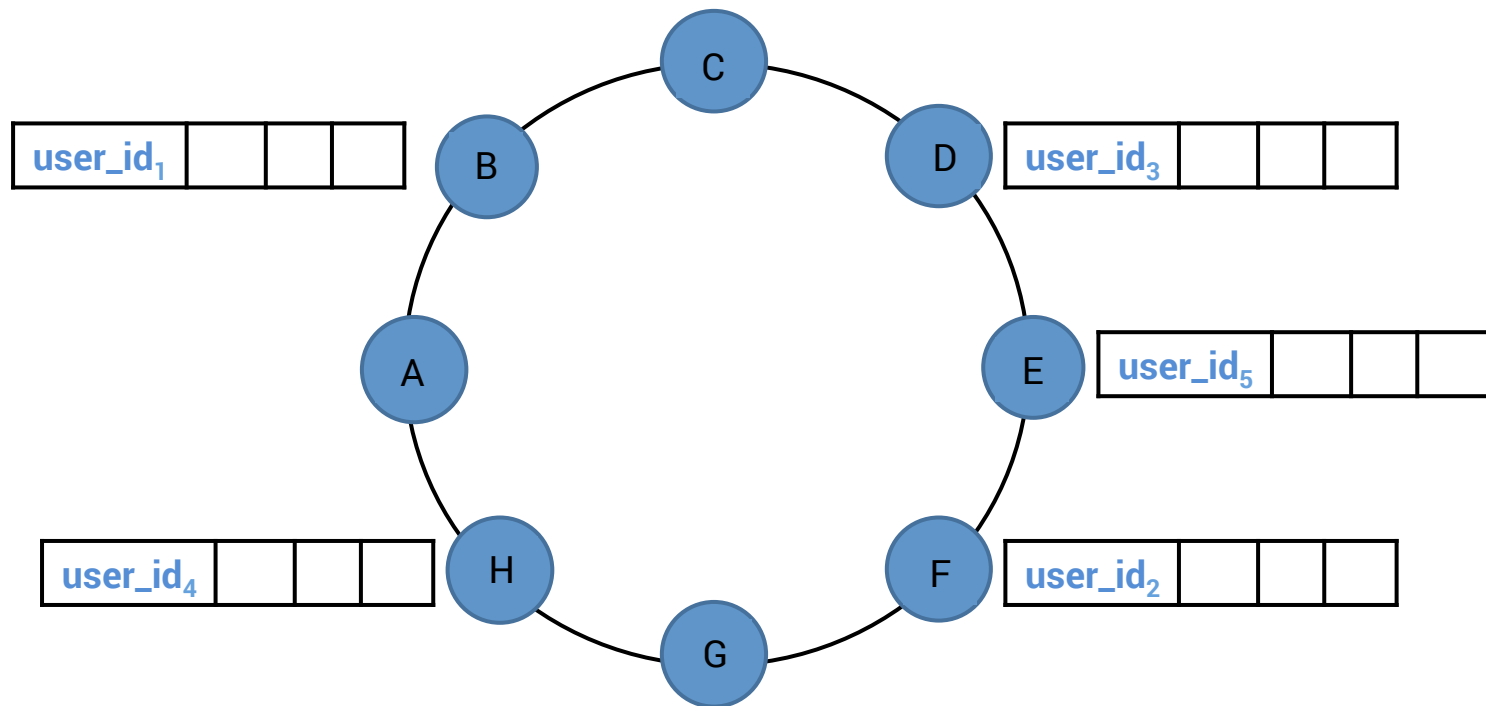
H: 7 /8,



# Distributed Table



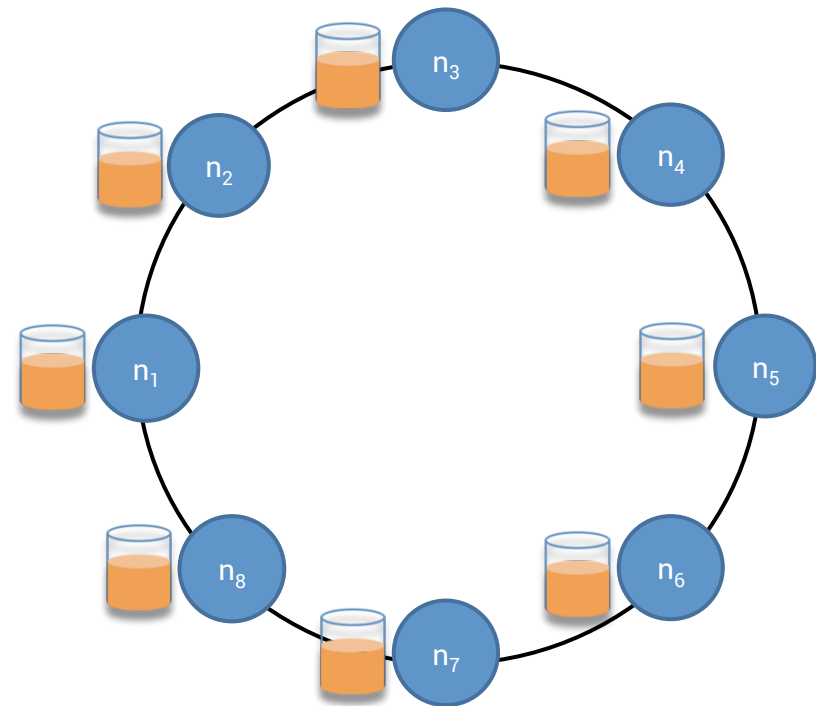
# Distributed Table



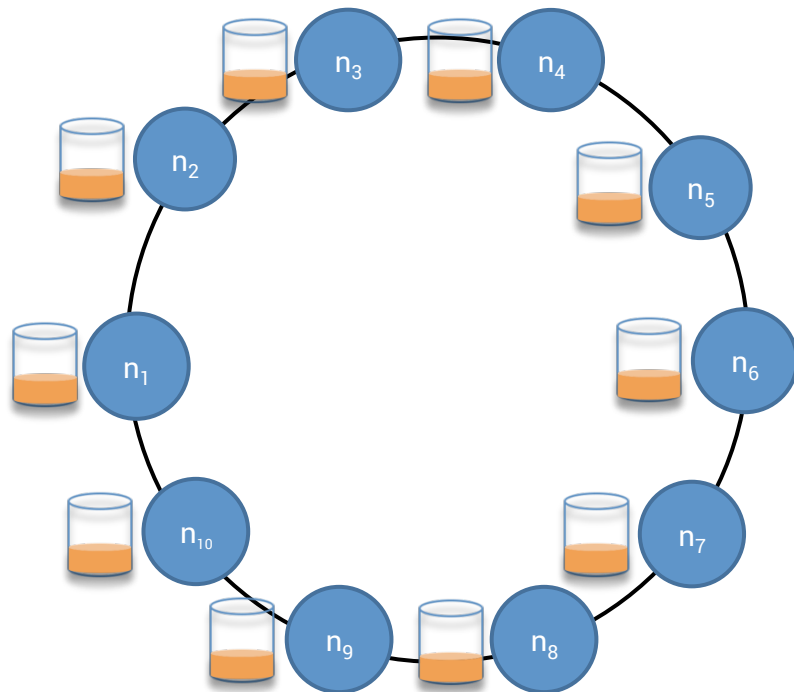


# Linear scalability

8 nodes

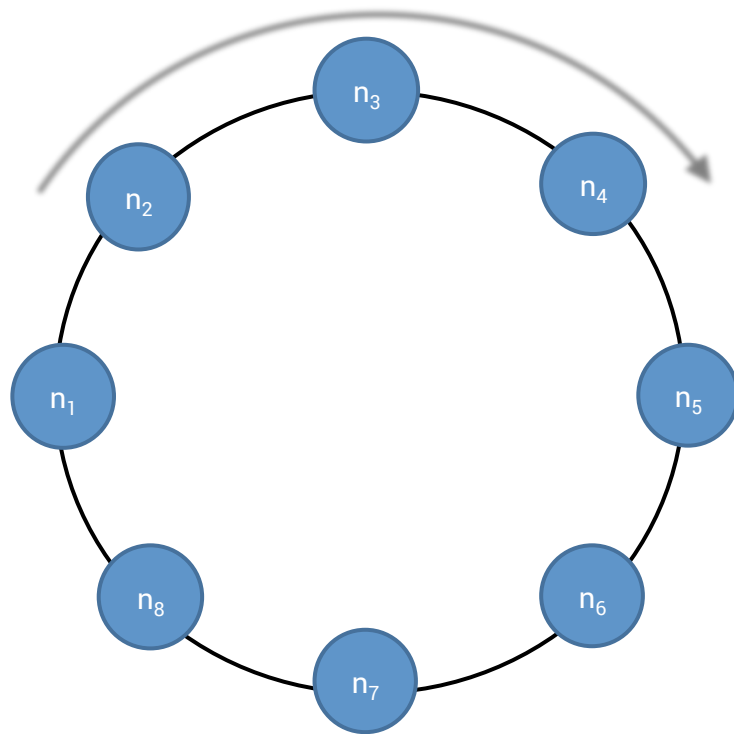


10 nodes



# Failure tolerance

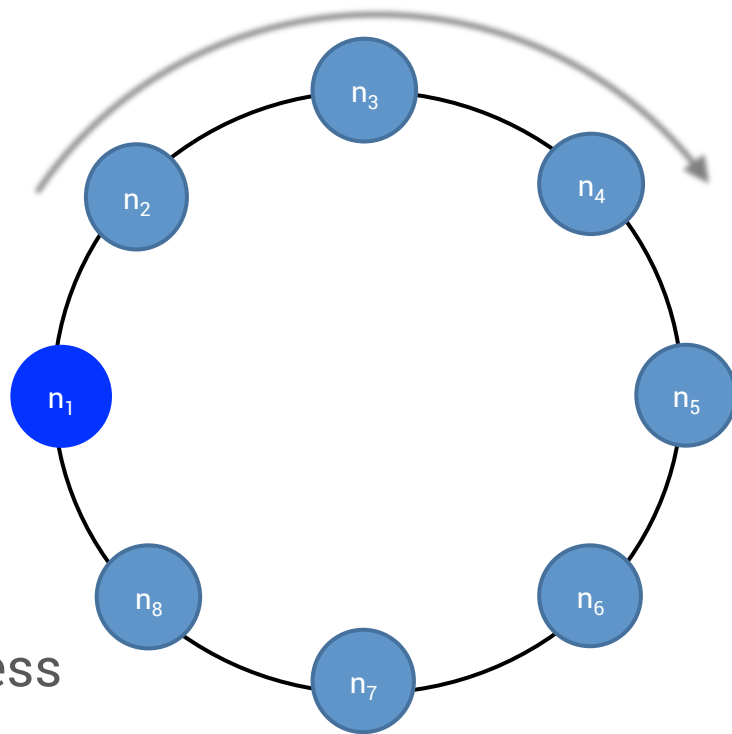
$$(RF) = 3$$



# Coordinator node

( / )

Coordinator



coordinator  $\rightarrow$  masterless

# Consistency

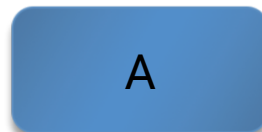
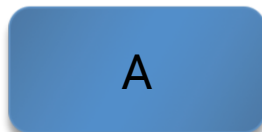
- ONE
- QUORUM (strict majority . . . RF)
- ALL

read & write

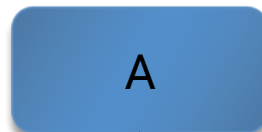
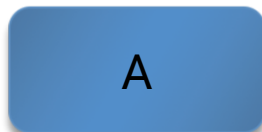
# Consistency in action

= 3, ONE, ONE

Write ONE: B



:



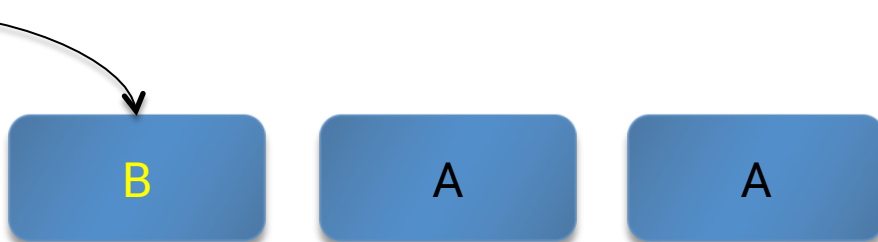
Read ONE: A



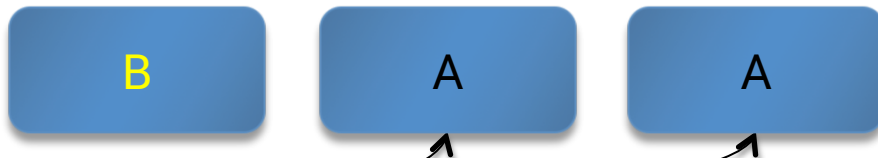
# Consistency in action

= 3, ONE, QUORUM

Write ONE: B



:

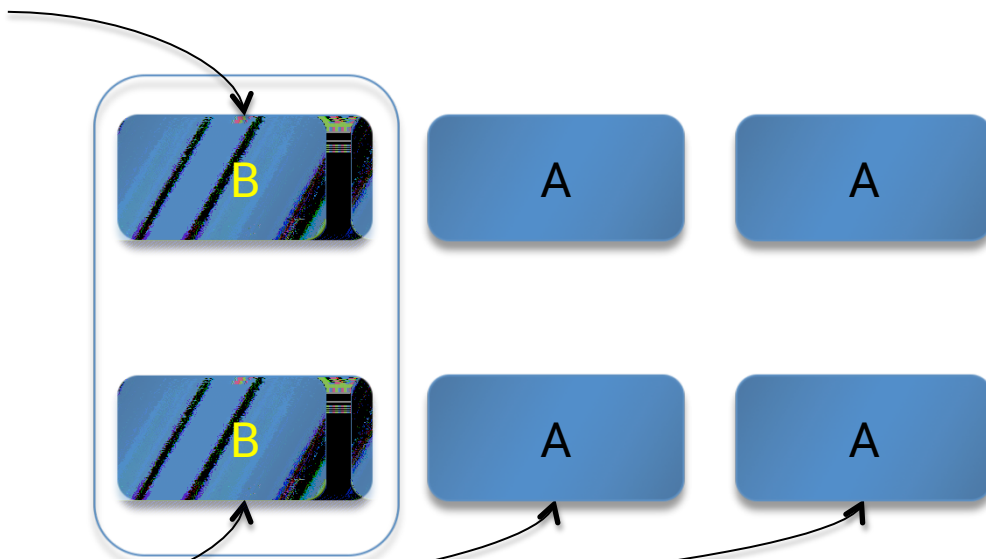


Read QUORUM: A

# Consistency in action

= 3, ONE, ALL

Write ONE: B

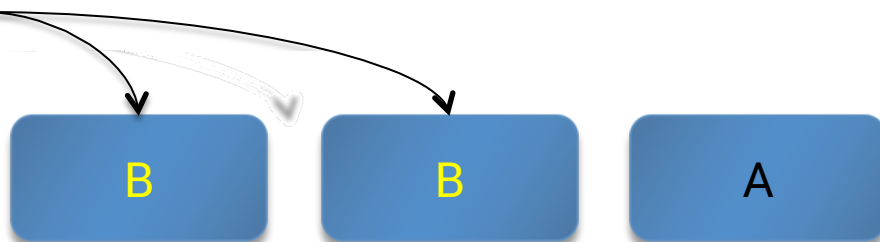


Read ALL: B

# Consistency in action

= 3, QUORUM, ONE

Write QUORUM: B



:



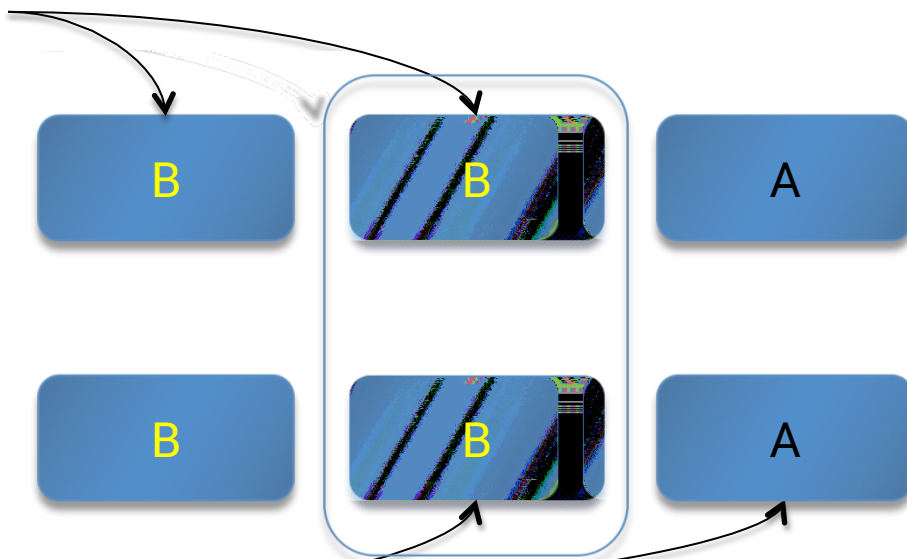
Read ONE: A



# Consistency in action

= 3, QUORUM, QUORUM

Write QUORUM: B



Read QUORUM: B



# Consistency trade-off

**Latency**

**Consistency**



# ONE

Fast, may not read latest written value

## QUORUM

Strict majority w.r.t. **Replication Factor**

Good balance

# Consistency level

# ALL

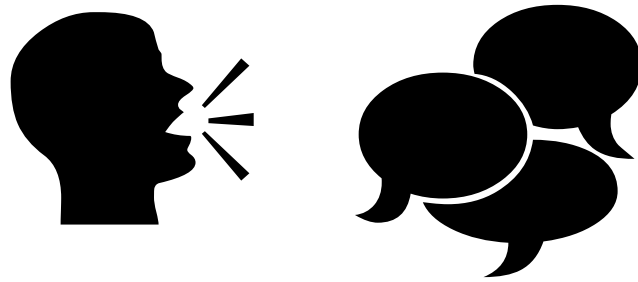
Paranoid

Slow, no high availability

# Consistency summary

**ONE**<sub>Read</sub> + **ONE**<sub>Write</sub>  
👉 available / (N-1)

**QUORUM**<sub>Read</sub> + **QUORUM**<sub>Write</sub>  
👉 available / 1+



Q & R

# Data model

Last Write Win

CQL basics

Clustered tables

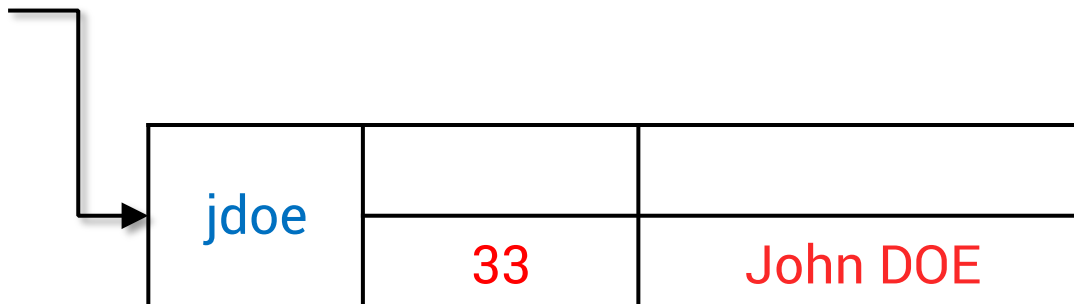
Lightweight transactions



# Last Write Win (LWW)

(login, , ) (jdoe, John DOE, 33);

#partition



jdoe		
jdoe	33	John DOE

# Last Write Win (LWW)

(login, , ) (jdoe, John DOE, 33);

auto-generated timestamp

The diagram illustrates the LWW (Last Write Win) concept. An 'auto-generated timestamp' is shown as a vertical line with a horizontal branch that splits into two arrows, each pointing to a row in a table. The table has three columns. The first column contains the value 'jdoe'. The second and third columns contain values from two different writes, both associated with the same timestamp '(t1)'. The first row has '(t1)' in the second column and '33' in the third column. The second row has '(t1)' in the second column and 'John DOE' in the third column.

jdoe	(t <sub>1</sub> )	33
	(t <sub>1</sub> )	John DOE

# Last Write Win (LWW)

= 34      login = jdoe;

SSTable <sub>1</sub>		
jdoe	(t <sub>1</sub> )	(t <sub>1</sub> )
	33	John DOE

SSTable <sub>2</sub>	
jdoe	(t <sub>2</sub> )
	34

# Last Write Win (LWW)

login = jdoe;

SSTable <sub>1</sub>		
jdoe	(t <sub>1</sub> )	(t <sub>1</sub> )
	33	John DOE

SSTable <sub>2</sub>	
jdoe	(t <sub>2</sub> )
	34

SSTable <sub>3</sub>	
jdoe	(t <sub>3</sub> )
	☒

tombstone

# Last Write Win (LWW)

login = jdoe;

?

SSTable<sub>1</sub>

jdoe	(t <sub>1</sub> )	(t <sub>1</sub> )
	33	John DOE


?

SSTable<sub>2</sub>

jdoe	(t <sub>2</sub> )	
	34	

?

SSTable<sub>3</sub>

jdoe	(t <sub>3</sub> )	
		

# Last Write Win (LWW)

`login = jdoe;`

✗

SSTable<sub>1</sub>

jdoe	(t <sub>1</sub> )	(t <sub>1</sub> )
	33	John DOE


✗

SSTable<sub>2</sub>

jdoe	(t <sub>2</sub> )	
	34	

✓

SSTable<sub>3</sub>

jdoe	(t <sub>3</sub> )	
		

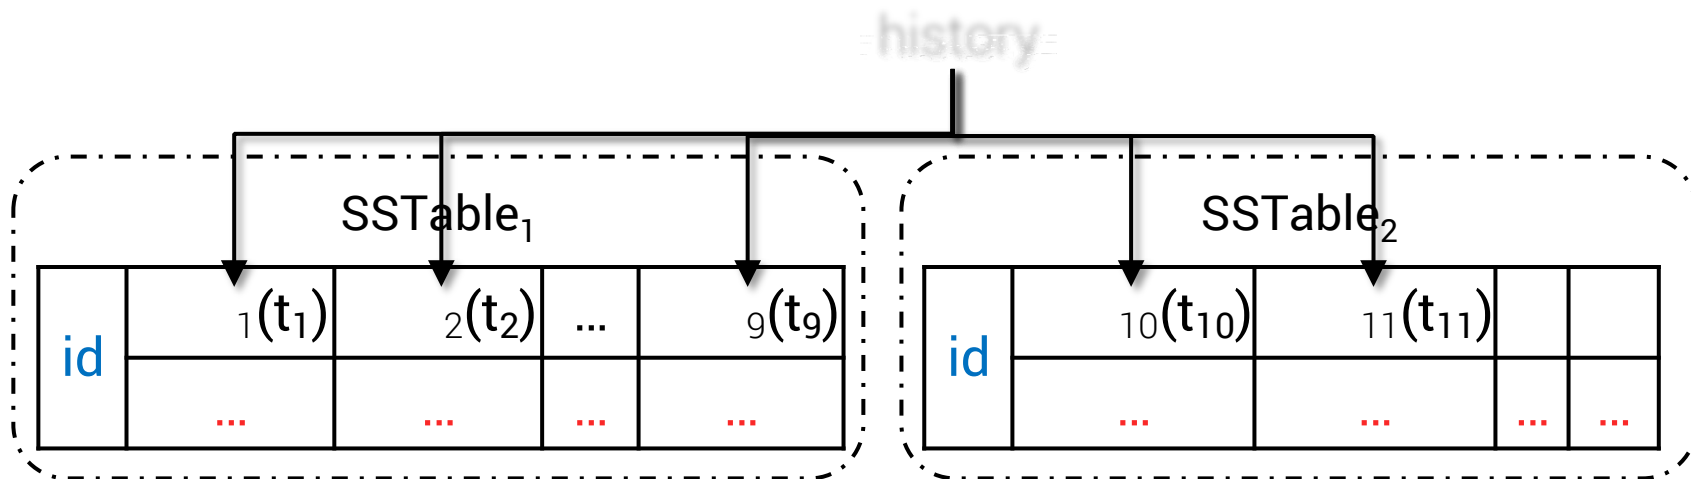
# Compaction



# Historical data

?

- do not
-  time-series





# CRUD operations

```
(login, , ) (jdoe, John DOE, 33);
```

```
= 34 login = jdoe;
```

```
login = jdoe;
```

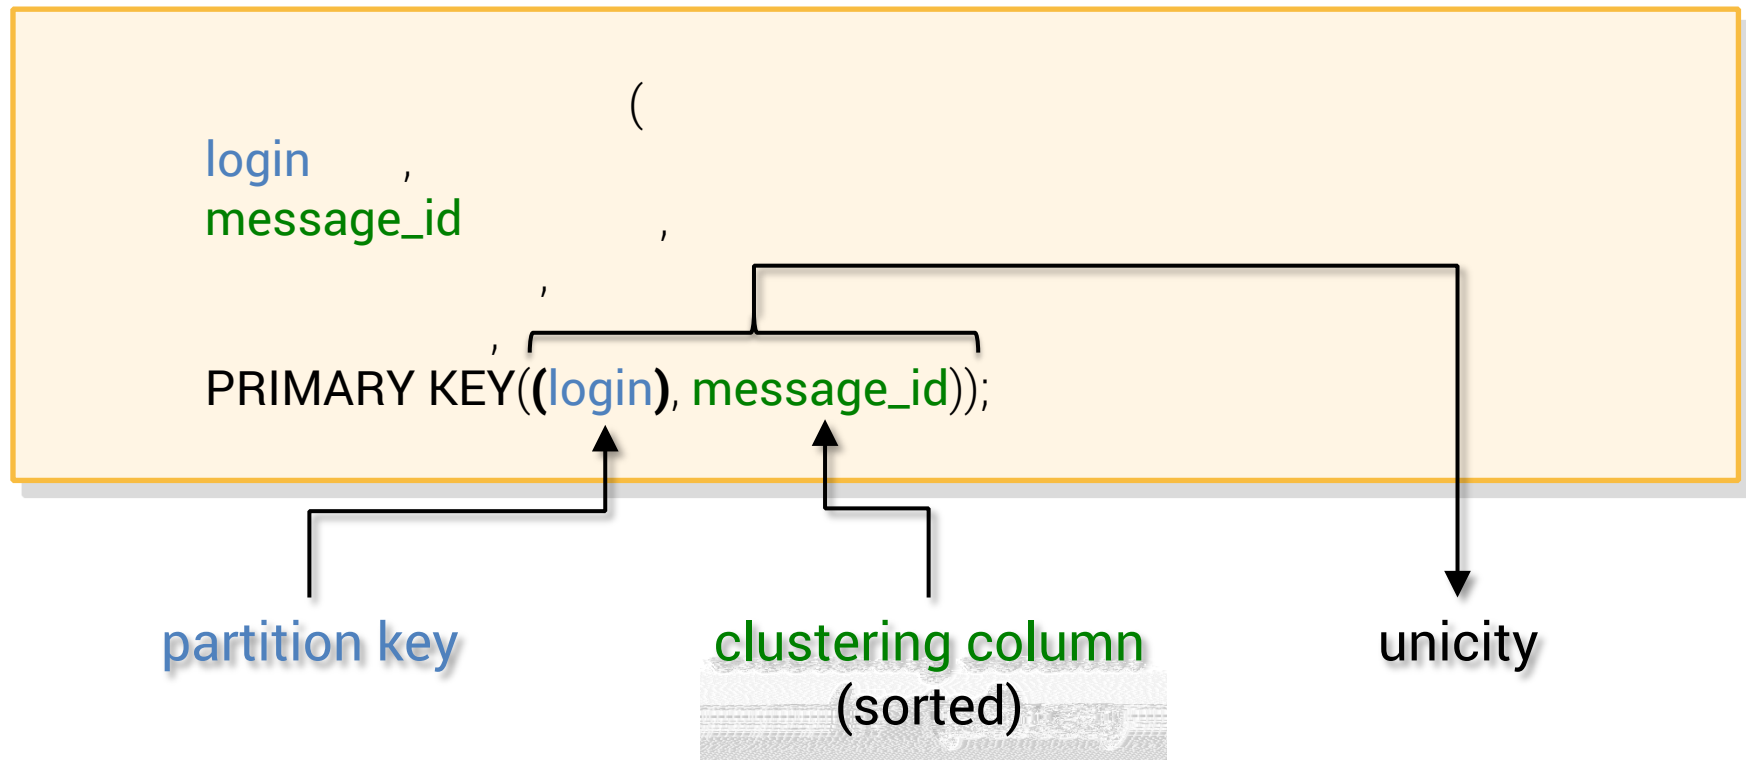
```
login = jdoe;
```

# Simple Table

login

(

# Clustered table (1 – N)



( )

```
* login = jdoe  
message_id = 2014-09-25 16:00:00 ;
```

```
* login = jdoe  
message_id <= 2014-09-25 16:00:00  
message_id >= 2014-09-20 16:00:00 ;
```

(#partition )

\*

message\_id = 2014-09-25 16:00:00 ;



(#partition )

\*

message\_id <= 2014-09-25 16:00:00  
message\_id >= 2014-09-20 16:00:00 ;



( #partition)

\*

login >= hsue

login <= jdoe;



( #partition)

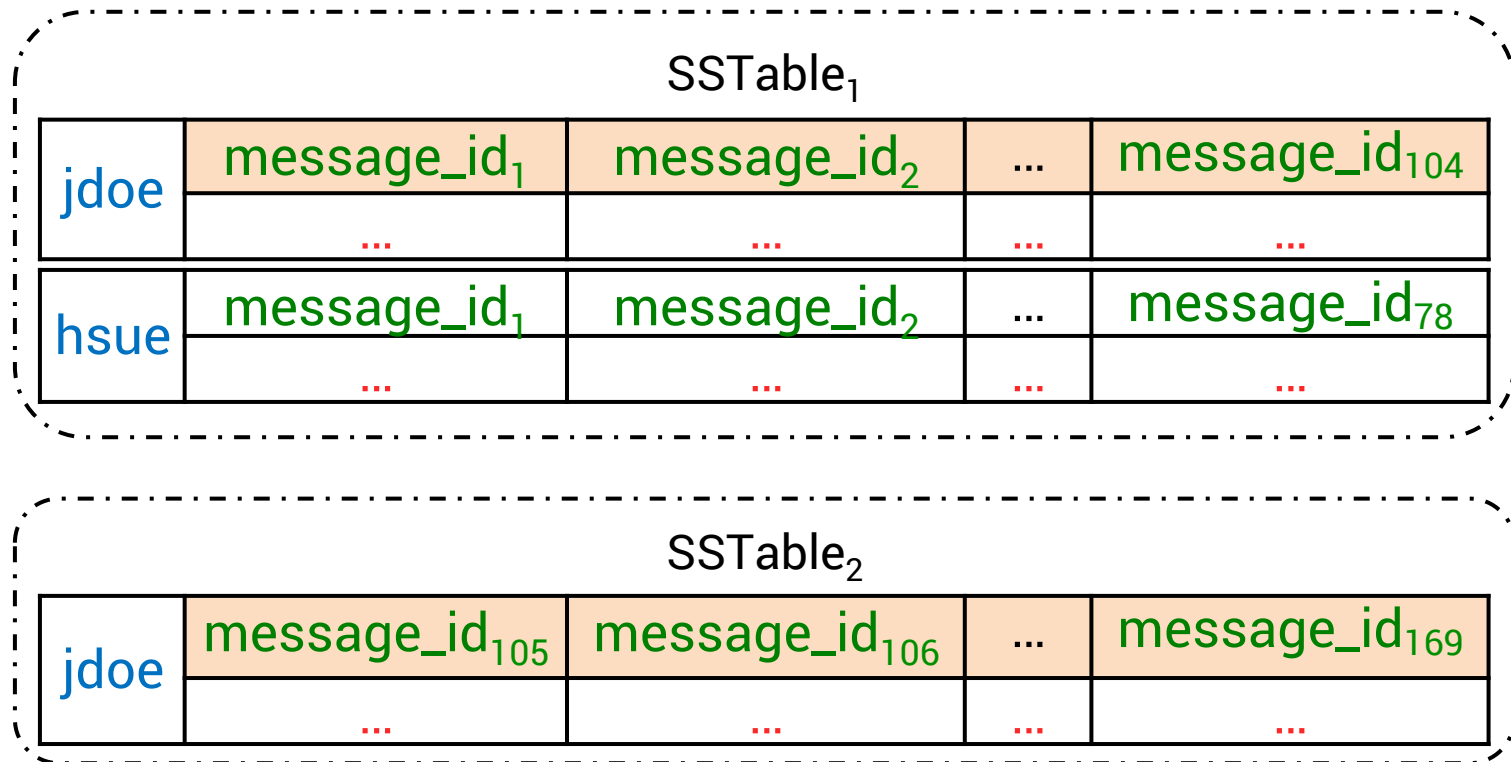
\*

login

%doe%;



# On disk layout



# Clustering order

```
(  
  login ,  
  message_id ,  
  ,  
  ,  
  ((login), message_id))  
CLUSTERING ORDER BY (message_id ) ;
```



# Reverse on disk layout

SSTable<sub>1</sub>

jdoe	message_id <sub>169</sub>	message_id <sub>168</sub>	...	message_id <sub>105</sub>
	...	...	...	...

SSTable<sub>2</sub>

jdoe	message_id <sub>104</sub>	message_id <sub>103</sub>	...	message_id <sub>1</sub>
	...	...	...	...

# WHERE clause restrictions

( / / / )

#partition

exact match (=) #partition,

(<, ≤, >, ≥)

-  full cluster scan

clustering columns,

(<, ≤, >, ≥)

exact match

WHERE

PRIMARY KEY

# Dynamic search

?

• ,

# Dynamic search

?

- ,

👉 Apache Solr ( ) (Datastax Enterprise)

```
* solr_query = age:[33 TO *] AND gender:male ;
```

```
* solr_query = lastname:*schwei?er ;
```

# Collections & maps

```
login (
  ,
  ,
  ,
  set<text>,
  list<text>,
  map<int, text>,
  : (
  PRIMARY KEY(login));
```

( $\approx 1000$ )

# From SQL to CQL

:

(

# From SQL to CQL

:

CQL **not** SQL

# From SQL to CQL

:

(

**no join**

(do you want to scale ?)



:

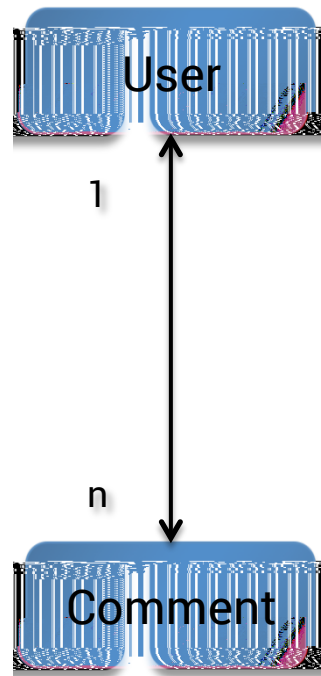
(

**no integrity constraint**

(do you want to read-before-write ?)

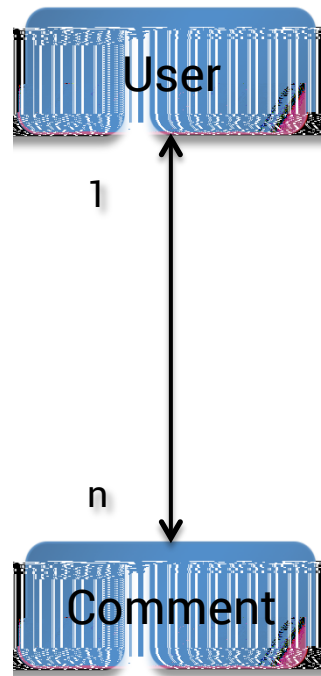
# From SQL to CQL

```
(  
  article_id  
  comment_id  
  author_id text, // typical join id  
  ((article_id), comment_id));
```



# From SQL to CQL

```
(  
  article_id  
  comment_id  
  author_json text, // de-normalize  
  ((article_id), comment_id));
```



# Data modeling best practices

- 
- $1 \approx 1$

# Data modeling best practices

- 
- $1 \approx 1$
- , necessary & immutable data
- / trade-off

# Data modeling best practices

## Article title

## Person JSON

- firstname/lastname
- date of birth
- gender
- mood
- location

John DOE

Male 33

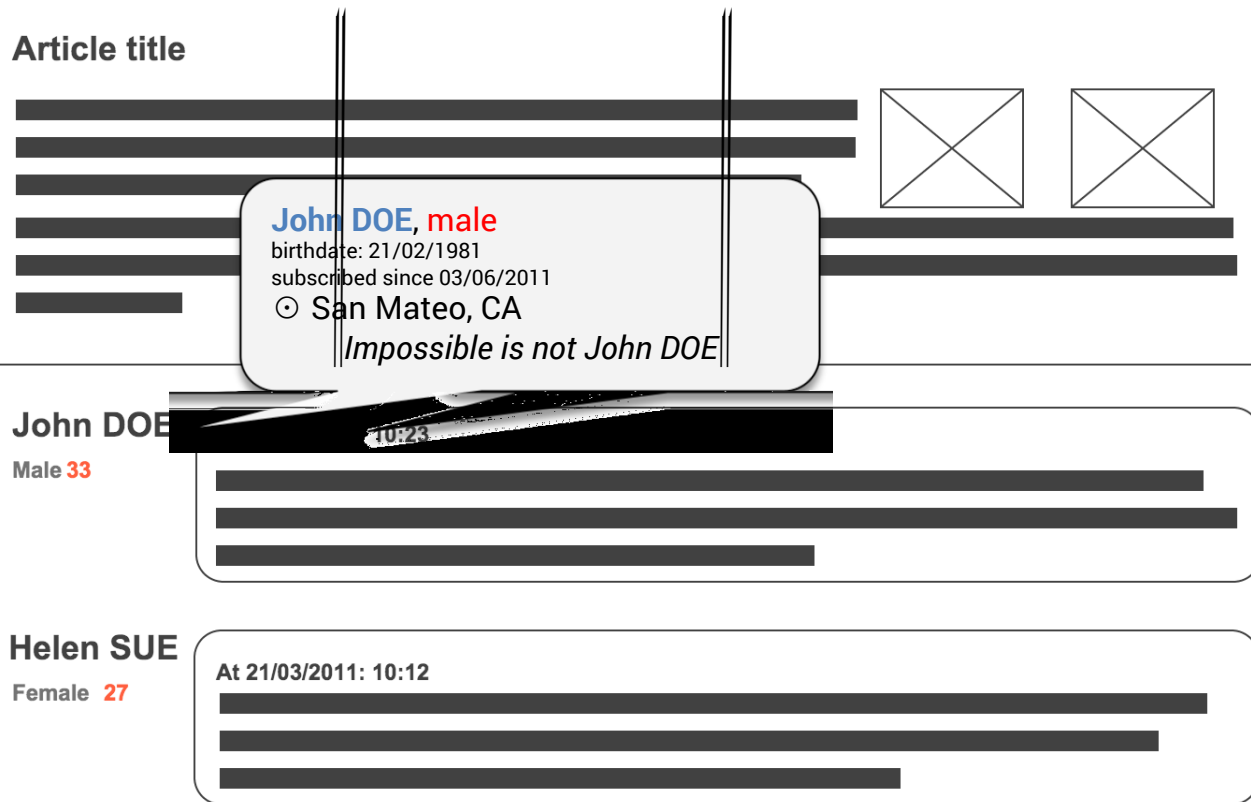
At 21/03/2011: 10:23

Helen SUE

Female 27

At 21/03/2011: 10:12

# Data modeling best practices



Full detail read from  
User table on click

# Lightweight Transaction (LWT)

? 

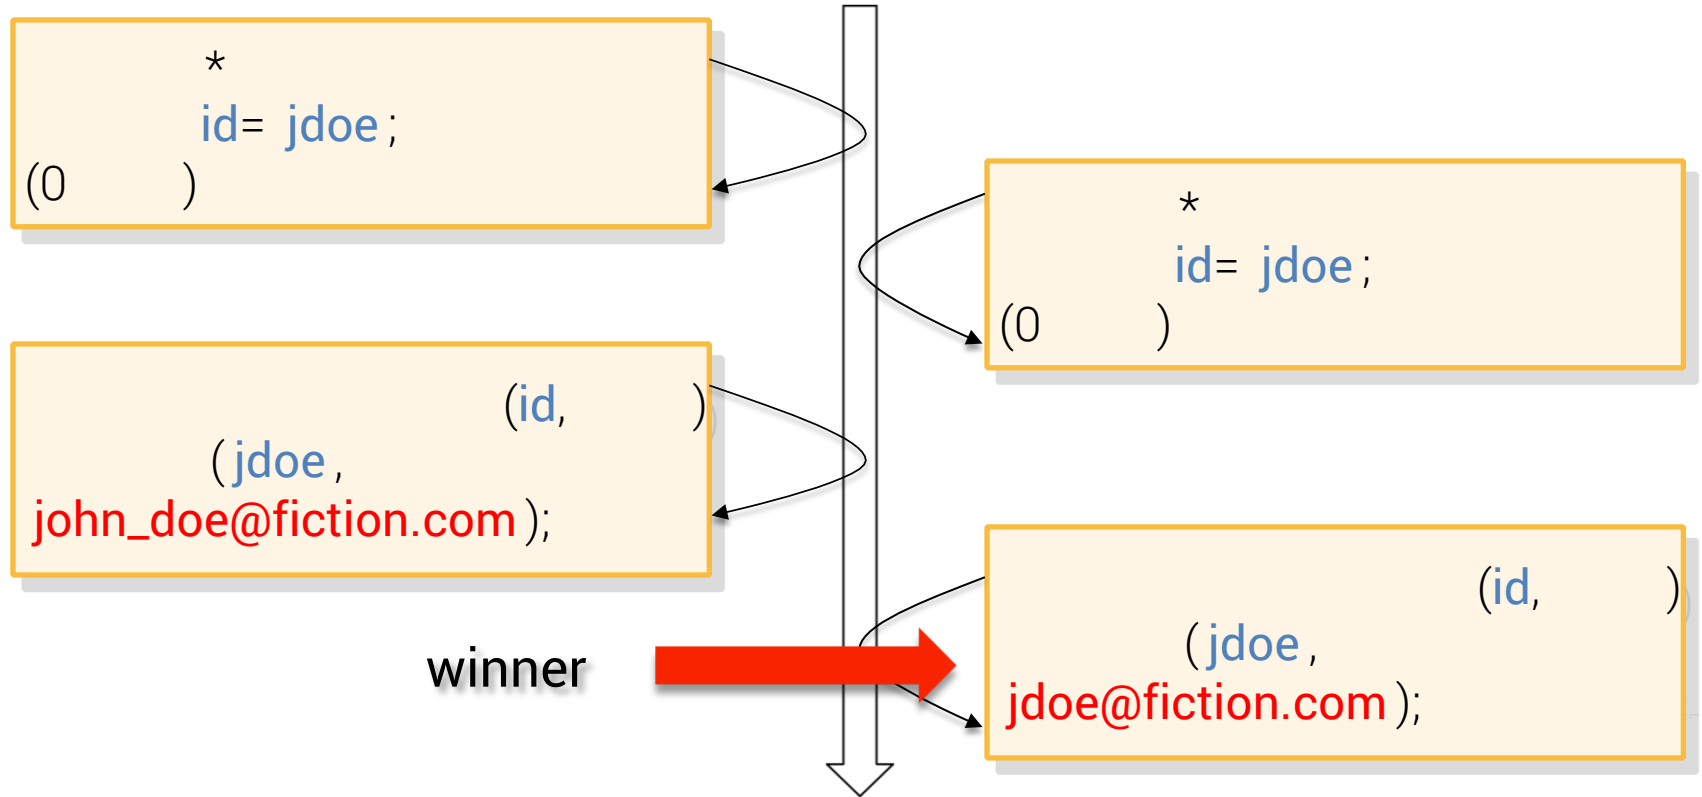
linearizable

? 

---



# Lightweight Transaction (LWT)



# Lightweight Transaction (LWT)

? 

Paxos

?

```
IF NOT EXISTS; ( , ) ( , . )  
=  
IF email = 'john_doe@fiction.com' = ;
```

# Lightweight Transaction (LWT)

-  must



IF NOT EXISTS  
IF EXISTS

:

:

(

(

# Lightweight Transaction (LWT)

-  \_\_\_\_\_  
must

= IF condition\_column = yyy

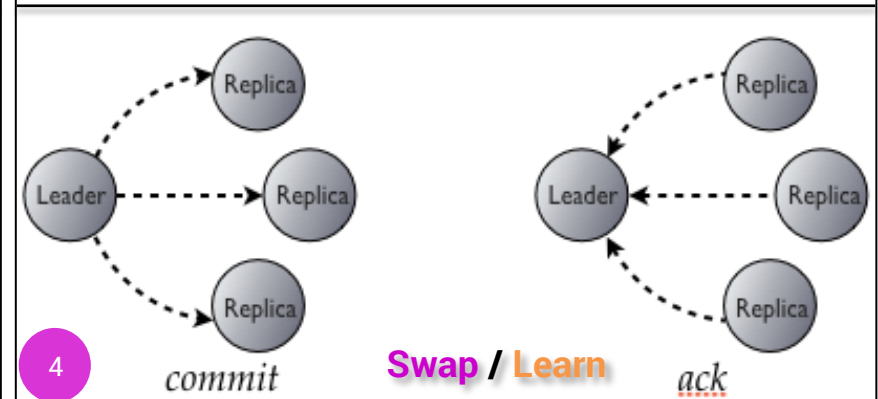
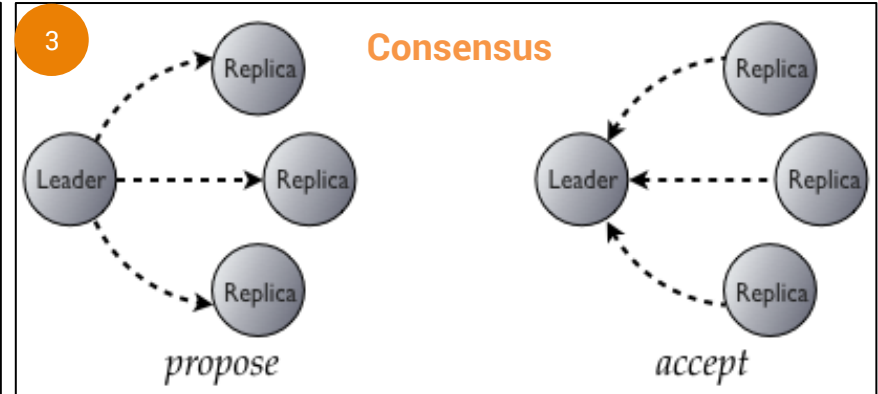
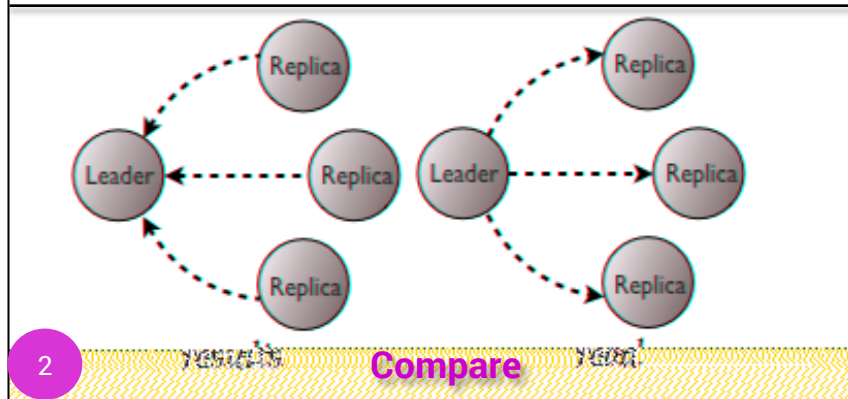
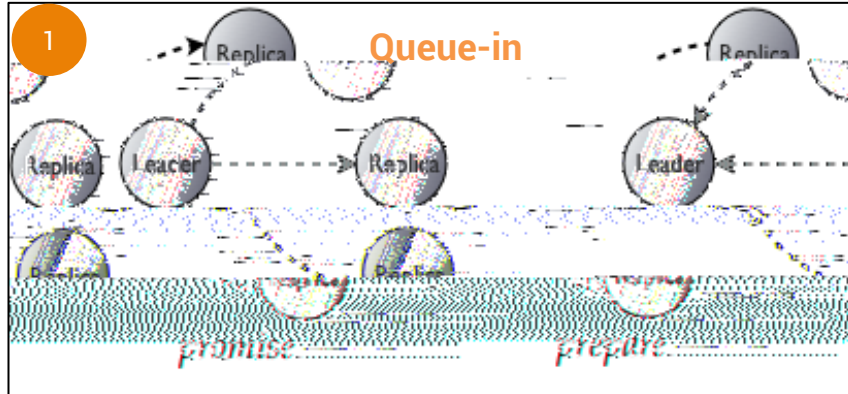


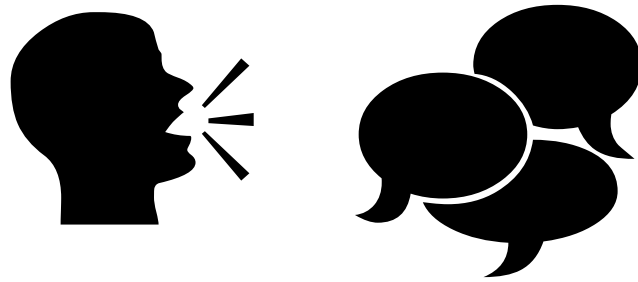
condition\_column = IF :

# Lightweight Transaction (LWT)

- (4 - ), **do not abuse**
- 1% – 5%

# Lightweight Transaction (LWT)





Q & R

# Thank You



**@doanduyhai**



**duy\_hai.doan@datastax.com**

**<https://academy.datastax.com/>**