

# KillrChat Exercises Handbook

---

DuyHai DOAN, Technical Advocate

# KillrChat presentation

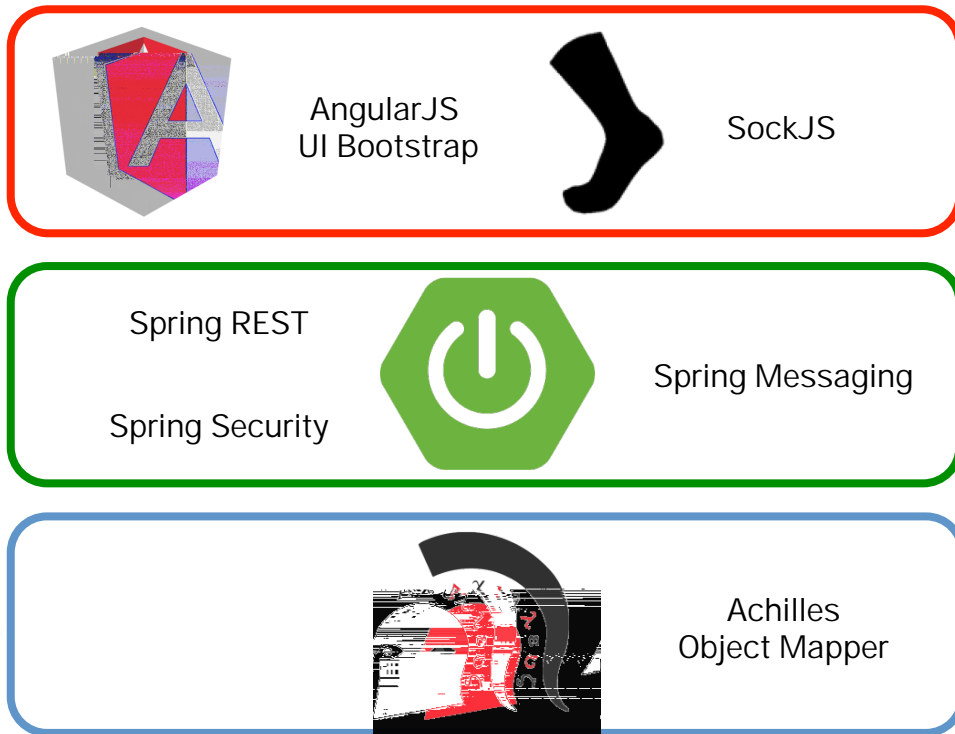

## What is KillrChat ?

- scalable messaging app

## Why KillrChat ?




- show real life de-normalization
- DIY exercise
- provide real application for attendees
- highlight Cassandra eco-system

# Technology stack




# Front end layout

KillrChat powered by Apache Cassandra™


  
My rooms

cassandra

killrchat



cassandra  
Cassandra room



rbu rbu joins the room

21:48:45

Hahahaha

21:48:49 DuyHai DOAN

yoooo

21:48:52 rbu RBU

oh yeah

21:48:59 Sébastien LE MERDY

plop

21:49:13 rbu RBU

<3

21:49:17 rbu RBU

J'ai déjà les web sockets

21:50:15 DuyHai DOAN

test

21:50:20 Sébastien LE MERDY

Alban Phelip joins the room

21:53:18

Ouai ouai

21:53:31 Alban PHELIP

New message

:D

Participants

Alban PHELIP

DuyHai DOAN

rbu RBU

Sébastien LE MERDY

# Exercises outline

TDD style

Implement the services to make tests green

Glue-code and front-end code provided

# Getting started

Clone the Git repository

```
git clone https://github.com/doanduyhai/killrchat.git
```

Go into the '*killrchat*' folder and launch tests

```
cd killrchat  
mvn clean test
```

# Exercise 1

User account management

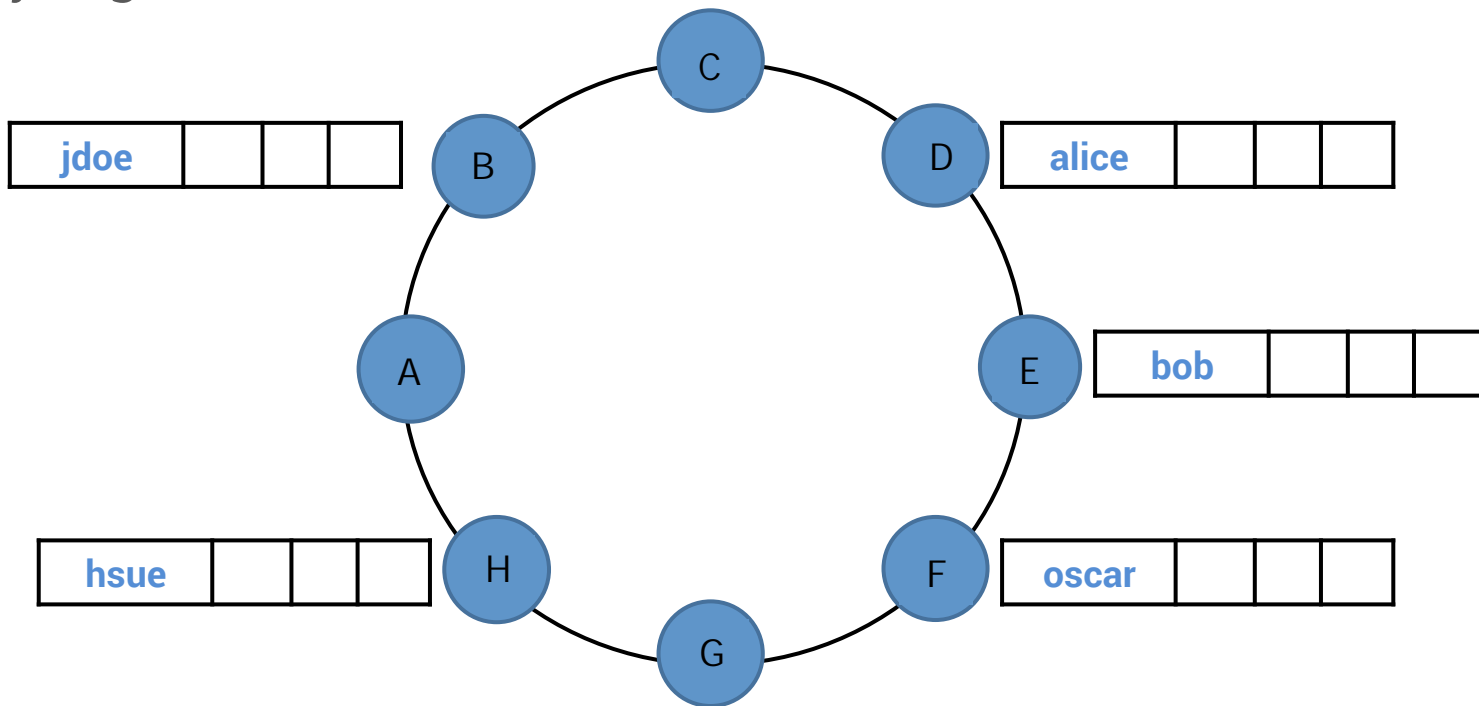
# Specifications

```
git checkout exercise_1_specs
```



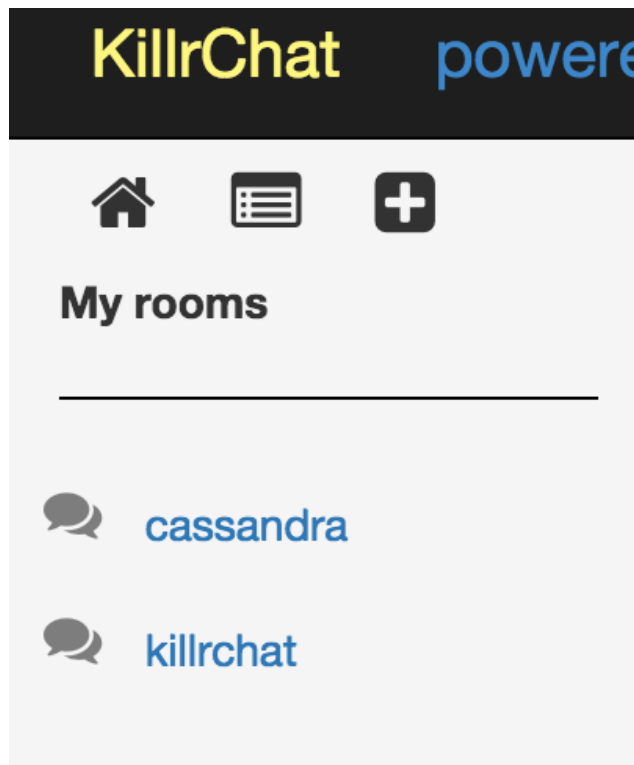
# Scalability

## Scaling by login



```
CREATE TABLE killrchat.users(  
    login text,  
    pass text, //password is not allowed because reserved word  
    lastname text,  
    firstname text,  
    bio text,  
    email text,  
    chat_rooms set<text>,  
    PRIMARY KEY(login));
```

# User's chat rooms



# User's chat rooms data model

How to store chat rooms for an user ?

```
CREATE TABLE killrchat.user_rooms(  
    login text,  
    room_name text,  
    PRIMARY KEY((login), room_name));
```

- pros: can store huge room count per user ( $10^6$ )
- cons: separated table, needs 1 extra SELECT

# User's chat rooms data model

Best choice

```
CREATE TABLE killrchat.users(  
    login text,  
    ...  
    chat_rooms set<text>, //list of chat rooms for this user  
    PRIMARY KEY(login));
```

- 1 SELECT fetches all data for a given user
- usually, 1 user is not in more that 1000 rooms at a time
- stores only room name

# Lightweight Transaction

Avoid creating the same login by 2 different users ?

👉 use Lightweight Transaction

```
INSERT INTO killrchat.users(room_name, ...)  
VALUES ('jdoe', ...) IF NOT EXISTS ;
```

Expensive operation

👉 do you create a new account every day ?

# Let's code!

## Tasks

- annotate UserEntity
- implement UserService

## Solution

```
git checkout exercise_1_solution
```

## Exercise 2


Chat room management



# Specifications

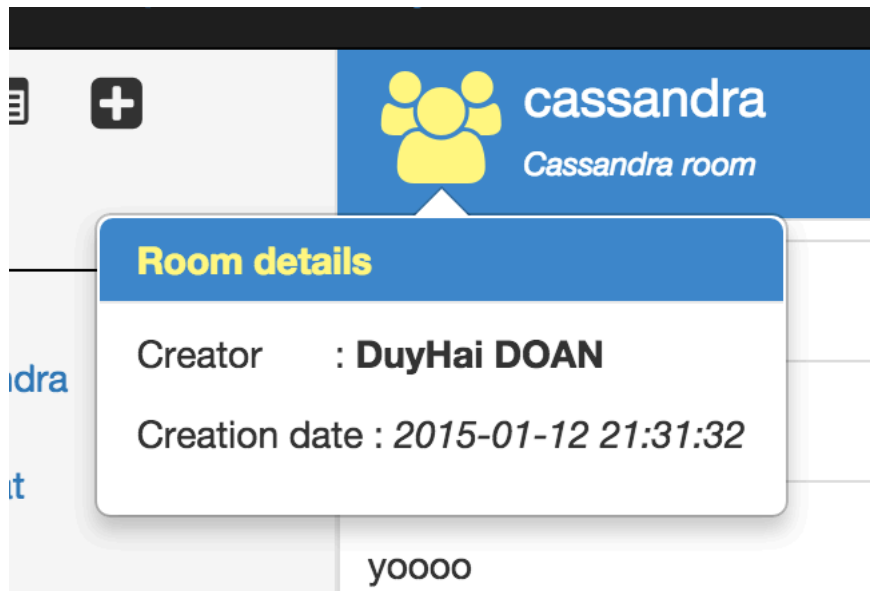
```
git checkout exercise_2_specs
```

# Scalability

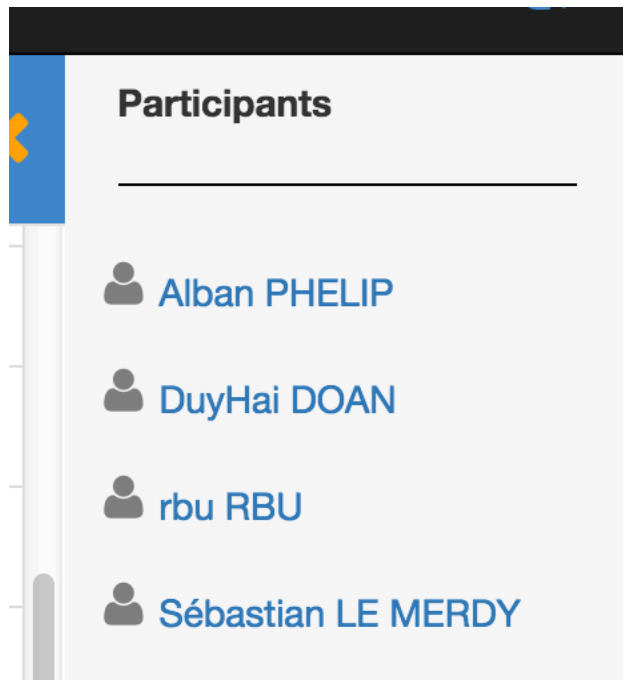


```
CREATE TABLE killrchat.chat_rooms(  
    room_name text,  
    name text,                // same as room_name  
    creation_date timestamp,  
    banner text,  
    creator text,             // de-normalization  
    creator_login text,  
    participants set<text>,   // de-normalization  
    PRIMARY KEY(room_name));
```

# Room details



# Room participants



# De-normalization

```
CREATE TABLE killrchat.chat_rooms(  
    room_name text,  
    ...  
    creator text,           // JSON blob {login: ..., firstname: ..., lastname: ...}  
    ...  
    participants set<text>, // JSON blob {login: ..., firstname: ..., lastname: ...}  
    PRIMARY KEY(room_name));
```

# Lightweight Transaction

Avoid creating the same room by 2 different users ?

👉 use Lightweight Transaction

```
INSERT INTO killrchat.chat_rooms(room_name, ...)  
VALUES ('games', ...) IF NOT EXISTS ;
```

# Listing all rooms

How to list all existing rooms ?

- limit to first 100 rooms
- rooms ordered by their token (hash of `room_name`)

Full text search ?

- possible with '*gam\**' semantics
- Lucene integration otherwise (DSE)



# Let's code!

## Tasks

- ChatRoomEntity already given with proper annotations
- Implement first methods in ChatRoomService

## Solution

```
git checkout exercise_2_solution
```

## Exercise 3

Participants management

Room deletion

# Specifications

```
git checkout exercise_3_specs
```

# Participant joining

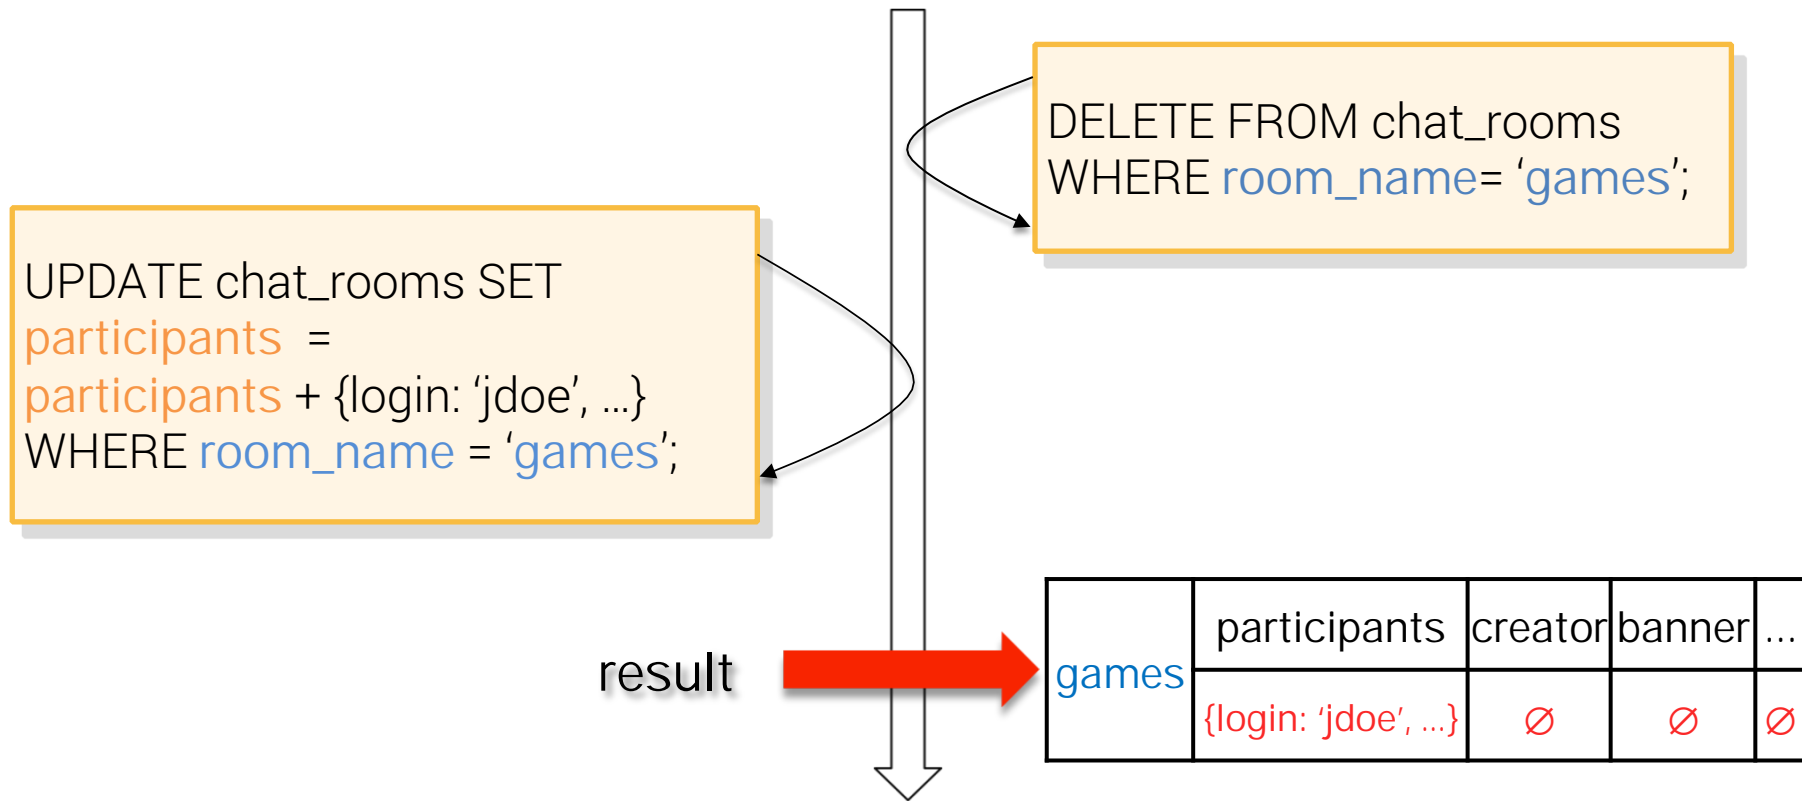
Adding new participant

```
UPDATE killrchat.chat_rooms SET participants = participants + {...}  
WHERE room_name = 'games';
```



What if the creator deletes the room at the same time ?

# Concurrent delete/update



# Participant joining

## Solution

☞ use Lightweight Transaction

```
UPDATE killrchat.chat_rooms SET participants = participants + {...}  
WHERE room_name = 'games' IF name = 'games';
```

Why not use UPDATE ... IF EXISTS ?

# Participant joining

## Solution

☞ use Lightweight Transaction

```
UPDATE killrchat.chat_rooms SET participants = participants + {...}  
WHERE room_name = 'games' IF name = 'games';
```

## Why not use UPDATE ... IF EXISTS ?

- syntax not yet available
- use column name = room\_name = partition key as condition
- see <https://issues.apache.org/jira/browse/CASSANDRA-8610>

# Participant leaving

Removing participant (no read-before-write)

```
UPDATE killrchat.chat_rooms SET participants = participants - {...}  
WHERE room_name = 'games';
```

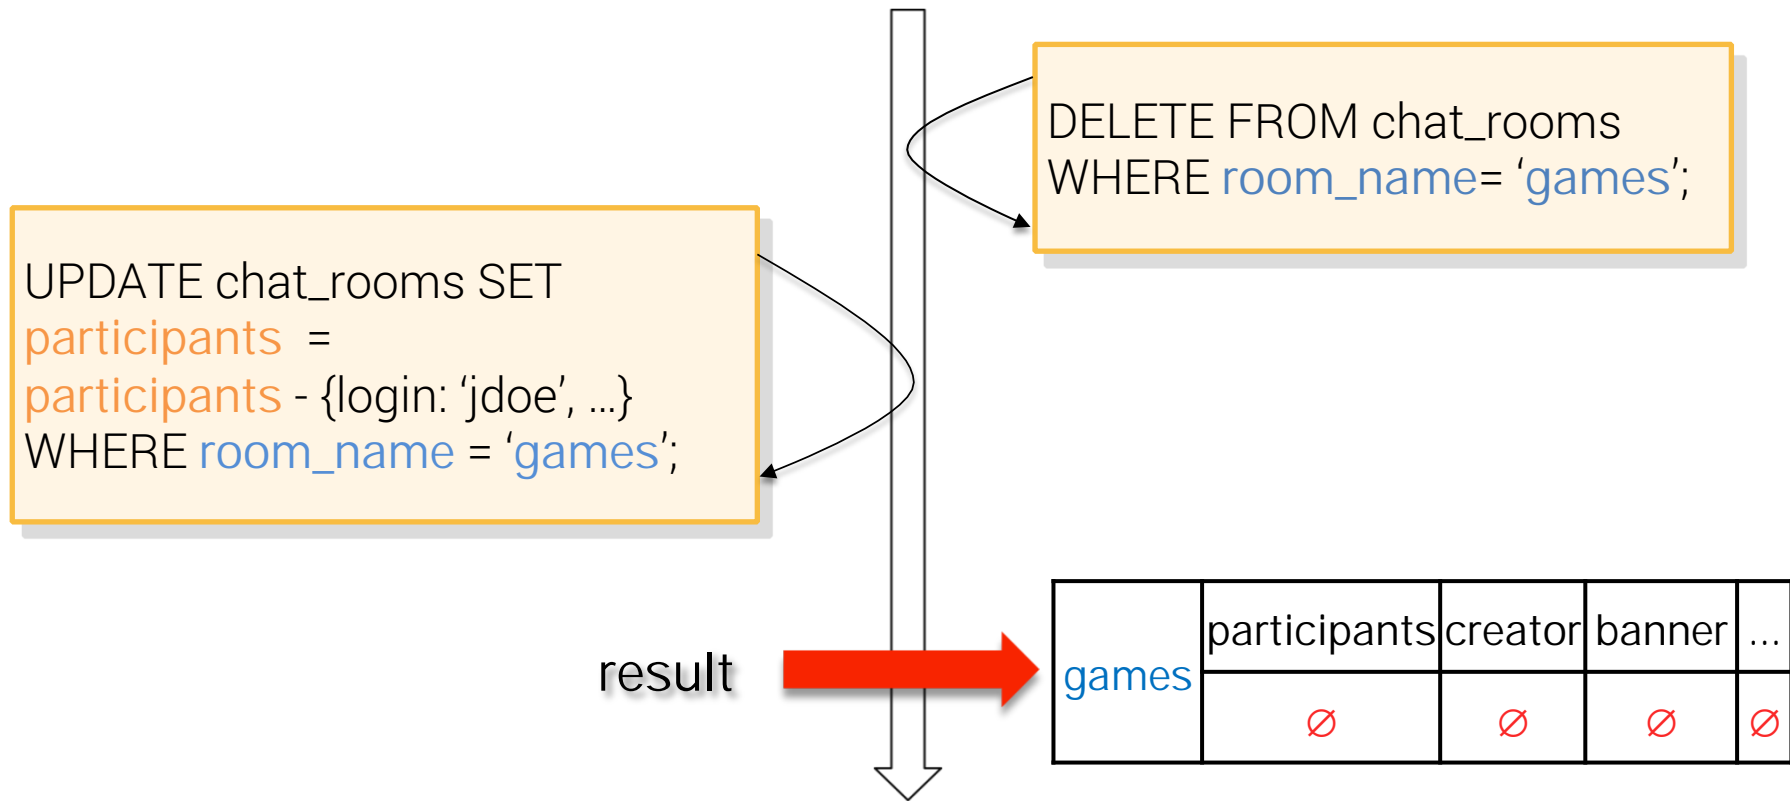


What if the creator deletes the room at the same time ?

- we'll create a tombstone
- tombstone will be garbage-collected by compaction



# Concurrent delete/update



# Deleting room

What if participant leaving at the same time ?

- not a problem, tombstone will be garbage

What if participant joining at the same time ?

👉 use Lightweight Transaction

Only room creator can delete room, no one else!

👉 use Lightweight Transaction

# Deleting room

## Solution

```
DELETE killrchat.chat_rooms  
WHERE room_name = 'games'  
IF creator_login = <current_user_login>;
```

## Advantages

- current user login coming from Security context, no cheating !
- slow but how often do you delete rooms ?

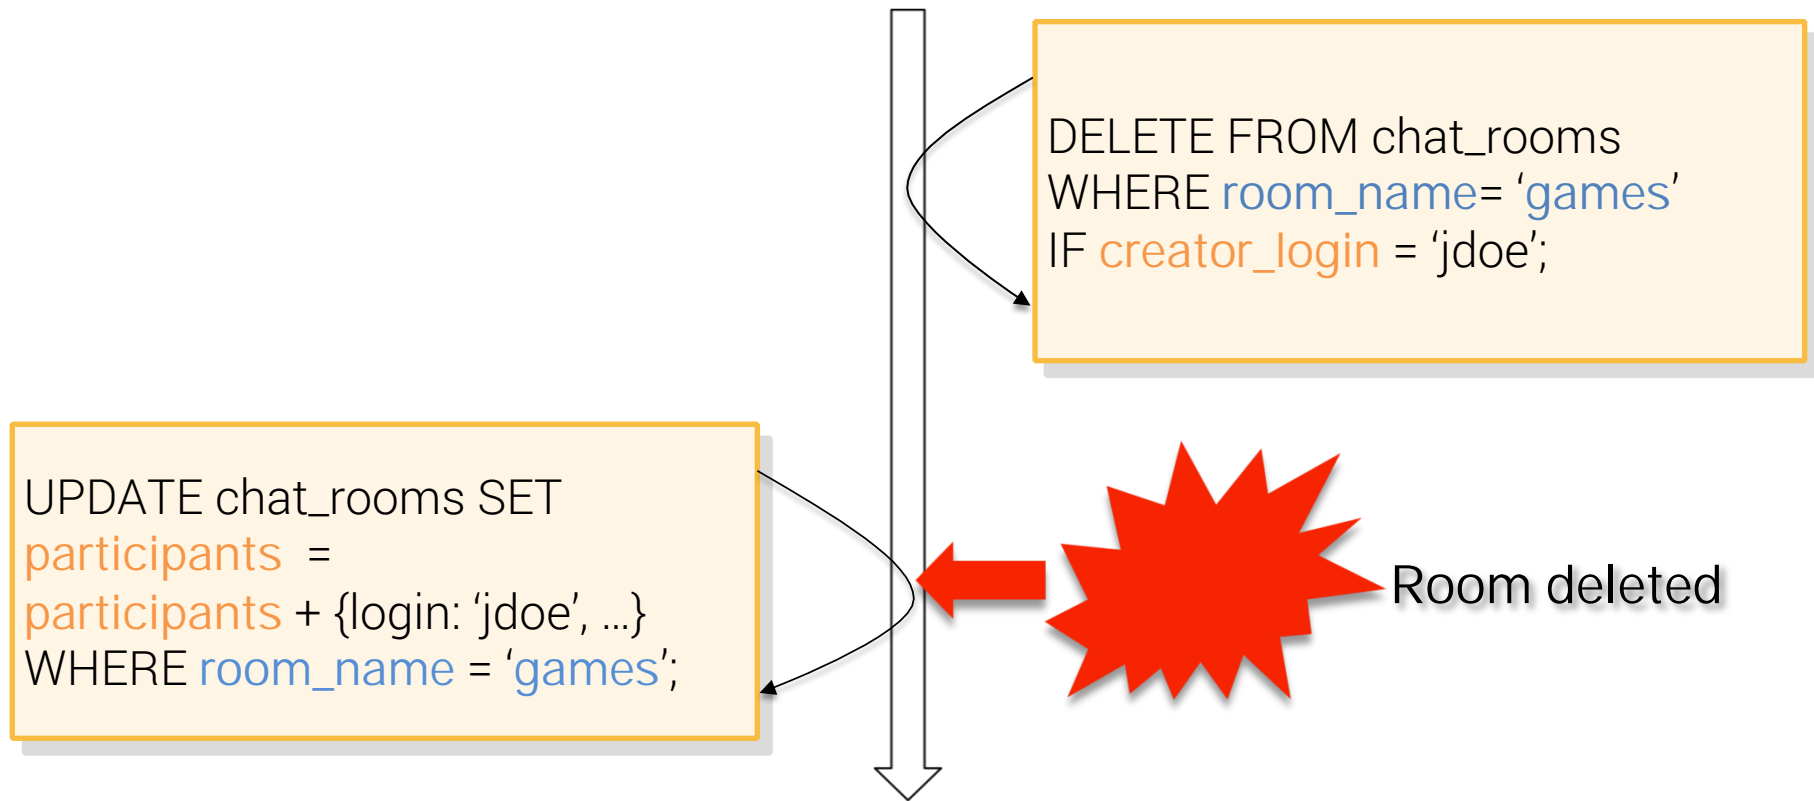
# Concurrent delete/update

```
UPDATE chat_rooms SET  
participants =  
participants + {login: 'jdoe', ...}  
WHERE room_name = 'games';
```

OK

```
DELETE FROM chat_rooms  
WHERE room_name = 'games'  
IF creator_login = 'jdoe';
```

# Concurrent delete/update



# Let's code!

## Tasks

- Implement remaining methods in ChatRoomService

## Solution

```
git checkout exercise_3_solution
```

## Exercise 4

Chat messages management


# Specifications

```
git checkout exercise_4_specs
```



# Scalability

Scaling messages by room name



```
CREATE TABLE killrchat.chat_room_messages(  
    room_name text,  
    message_id timeuuid,  
    content text,  
    author text,           // JSON blob {login: ..., firstname: ..., lastname: ...}  
    system_message boolean,  
    PRIMARY KEY((room_name), message_id)  
) WITH CLUSTERING ORDER BY (message_id DESC);
```

# Data model

Clustering column **message\_id** order by DESC

- latest messages first
- leverage the new row cache in Cassandra 2.1

## Improvements

- current data model limits messages count to  $\approx 500 \times 10^6$
- bucketing by day is the right design

```
PRIMARY KEY((room_name, day), message_id) //day format yyyyMMdd
```


# Let's code!

## Tasks

- MessageEntity already given with proper annotations
- Implement methods in MessageService

## Solution

```
git checkout exercise_4_solution
```



Q & R

# Thank You



**@doanduyhai**



**duy\_hai.doan@datastax.com**

**<https://academy.datastax.com/>**